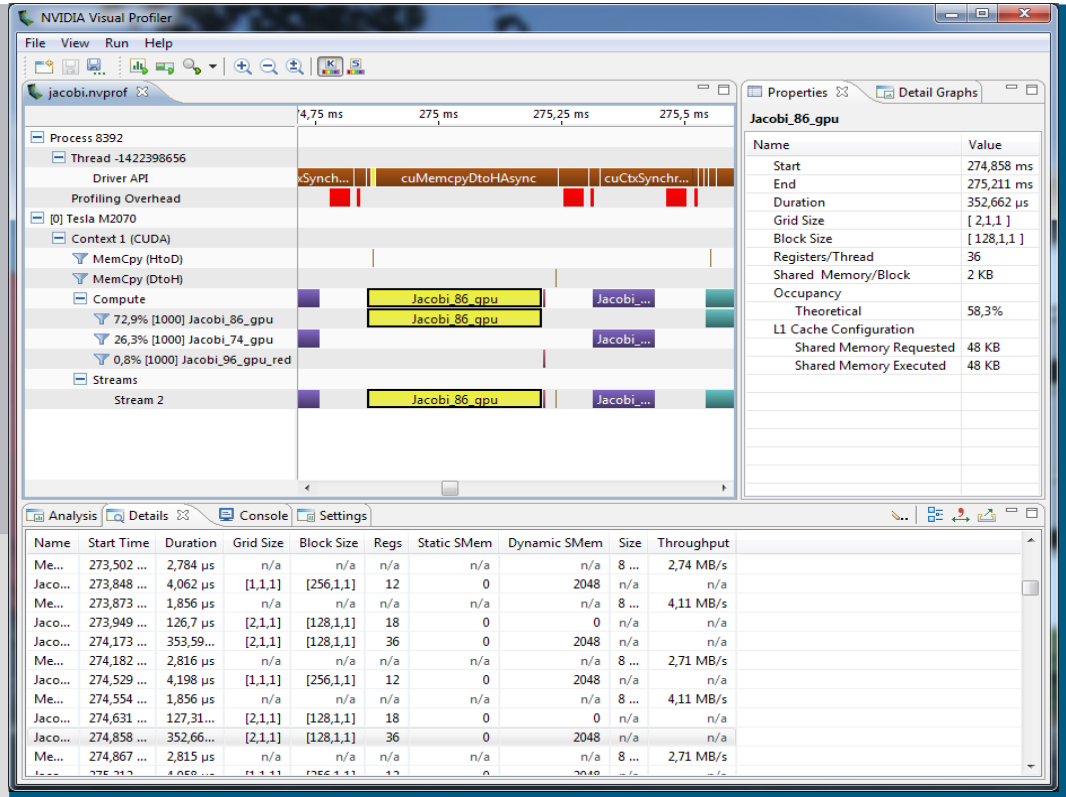

CUDA Performance Measurement

Jiri Kraus (NVIDIA)

NVIDIA

JÜLICH FORSCHUNGSZENTRUM

Why Performance Measurement Tools?

- You can only improve what you measure
  - Need to identify:
    - Hotspots: Which function takes most of the run time?
    - Bottlenecks: What limits the performance of the Hotspots?
- Manual timing is tedious and error prone
  - Possible for small application like jacobi and matrix multiplication
  - Impractical for larger/more complex application
- Access to hardware counters (PAPI, CUPTI)

The command line profiler nvprof

- Simple launcher to get profiles of your application
- Profiles CUDA Kernels and API calls

```
> nvprof ./jacobi
======== NVPROF is profiling jacobi...
======== Command: jacobi
Jacobi (serial)
[…] snip
======== Profiling result:
 Time(%)      Time   Calls       Avg       Min       Max  Name
   72.14  352.65ms    1000  352.65us  350.48us  354.94us  Jacobi_86_gpu
   26.02  127.23ms    1000  127.23us   93.48us  128.34us  Jacobi_74_gpu
    0.84    4.09ms    1000    4.09us    4.04us    4.36us  Jacobi_96_gpu_red
    0.61    3.00ms    1009    2.97us    2.78us   56.16us  [CUDA memcpy HtoD]
    0.39    1.91ms    1002    1.91us    1.82us   52.41us  [CUDA memcpy DtoH]
```

nvprof interoperability with nvvp

- nvprof can write the application profile to nvvp compatible file:

```
nvprof -o jacobi.nvprof ./jacobi
```

- Import in nvvp

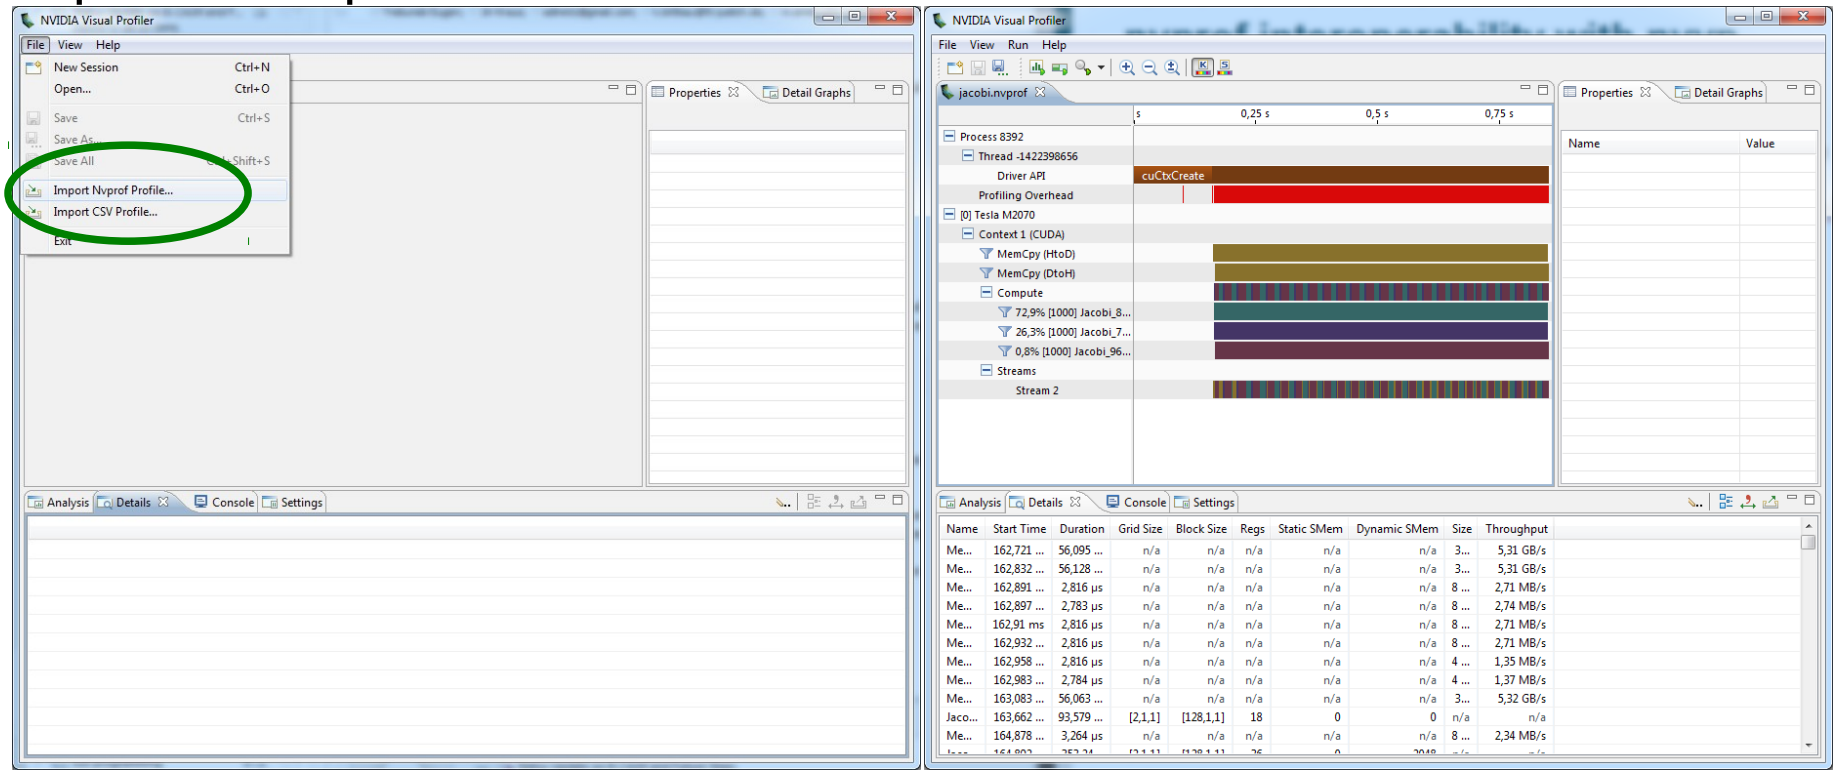

nvprof important command-line options

```
Options:
  -o,  --output-profile <filename>
                              Output the result file which can be imported later
                              or opened by the NVIDIA Visual Profiler.
 --events <event names>
                              Specify the events to be profiled on certain
                              device(s). Multiple event names separated by comma
                              can be specified. Which device(s) are profiled is
                              controlled by the '--devices' option. Otherwise
                              events will be collected on all devices.
                              For a list of available events, use
                              '--query-events'.
 --query-events               List all the events available on each device.
 -h,  --help                  Print this help information.
```

nvvp introduction

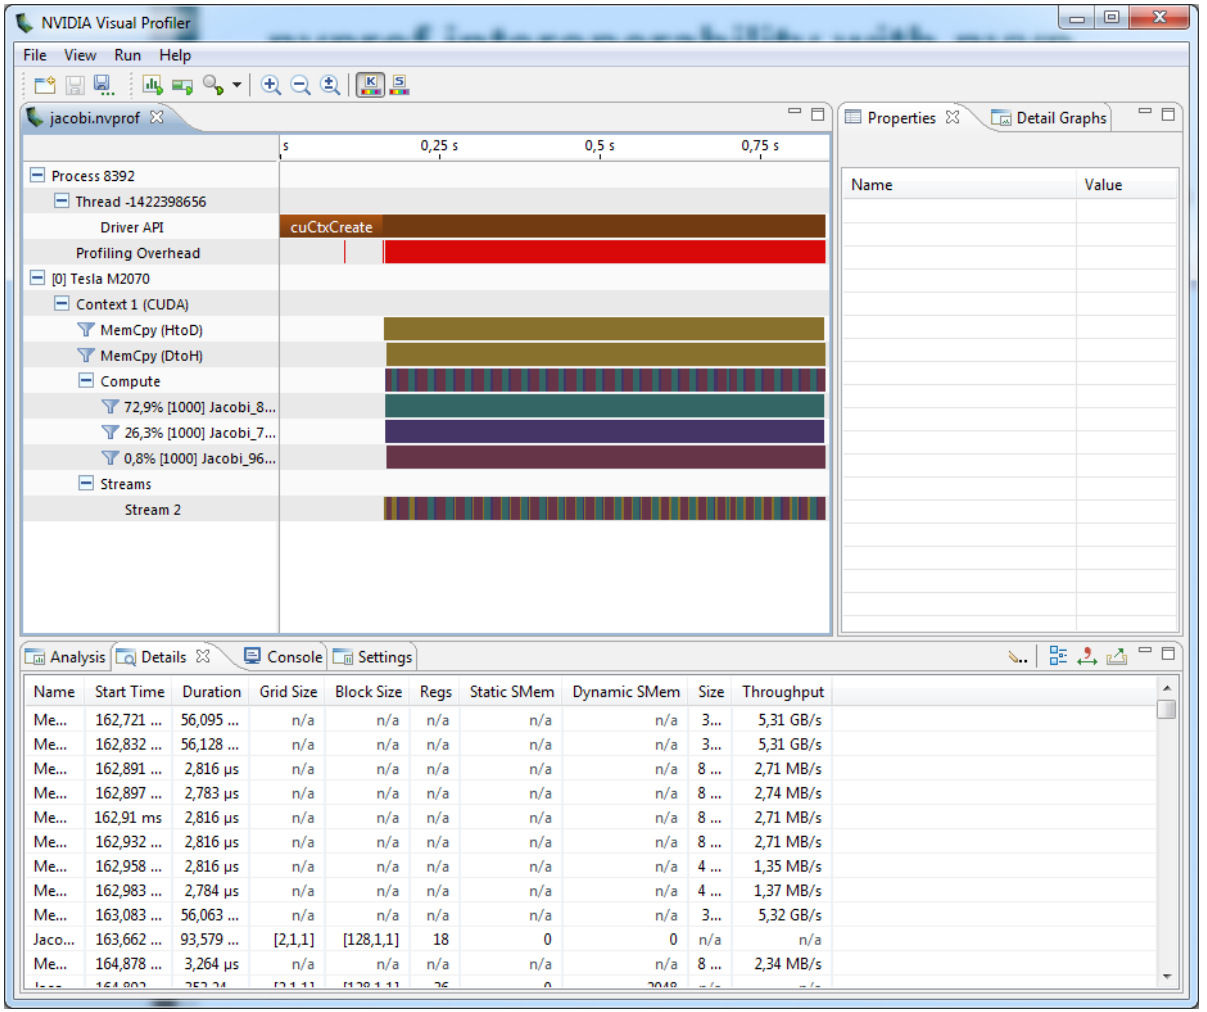

Task 1: Analyze Jacobi Timeline

- Start jacobi with nvprof and write profile to file
- Import profile into nvvp
- Compare the profiles with and without data region.

Task 2: Analyze matrix multiplication example with nvvp

- Start new session in nvvp with the matrix multiplication example
- Run the “Uncoalesced Global Memory“ experiment

Task 3: Analyze matrix multiplication example with nvprof

- Start matrix multiplication with nvprof and collect gld_inst_32bit event
- Import profile into nvvp
- Read the value of gld_inst_32bit and compare it to the size of the input matrices and the number of executed floating point multiplications

*Hint: M*N*K and M*N + N*K*

Cheat Sheet

- Start nvprof

```
nvprof -o <output-profile> ./a.out
```

- Start nvvp

```
nvvp
```

- profiler users guide

```
http://docs.nvidia.com/cuda/profiler-users-guide/index.html
```